

Abhinav Aditya

Bengaluru, India | abhinavaditya165@gmail.com | +91 7079020419 | abhinavaditya.com | linkedin.com/in/abhinav-aditya
github.com/abhinav162

PROFESSIONAL SUMMARY

Backend and applied AI engineer with 1.5+ years building production GenAI systems and cloud services end to end: multi-agent LLM pipelines, a multi-provider LLM gateway with automatic fallbacks, and messaging infrastructure dispatching 400K+ emails/day. Strong on scalability, reliability, and ownership across microservices, RabbitMQ, Redis, PostgreSQL, Docker, and AWS; moved from frontend intern to owning backend and AI systems within a year, and quick to pick up new stacks.

SKILLS

Backend & Messaging: Python (FastAPI), Node.js, Express, REST APIs, Microservices, RabbitMQ, AWS SQS, WebSockets, OAuth 2.0, JWT, Webhooks

AI & LLM Engineering: LLM tool-calling, multi-agent orchestration, RAG & embeddings, LLM gateway & model routing, MCP server design, OpenAI / Azure OpenAI / Gemini / Claude APIs, LangChain, streaming LLM UX, AI coding assistants (Claude Code, Cursor, Copilot)

Databases & Caching: PostgreSQL, MongoDB, Redis, Qdrant (vector search)

Cloud & DevOps: Docker, AWS (EC2, ECS, ECR, S3, Lambda), Terraform, Ansible, GitHub Actions CI/CD, Nginx, Grafana, Prometheus, Sentry

Languages: Python, TypeScript, JavaScript, SQL

Frontend: React, Next.js, Redux Toolkit, Tailwind CSS, Vitest / Playwright testing

WORK EXPERIENCE

HireQuotient | Software Development Engineer I

Bengaluru, India · Jan 2025 – Present

AI-powered recruiting and sourcing platform used by enterprise talent teams.

- Built and scaled a multi-provider outbound email engine (Gmail, Outlook, SendGrid, SMTP/IMAP) on **RabbitMQ and Redis**, dispatching **400K+ emails/day** and **1.5M+ in a peak week**, with sender-pool inbox rotation, retry handling, and OAuth 2.0 connection flows.
- Moved every OpenAI and Azure OpenAI call behind a single **LLM gateway**, so providers and models are configured rather than hard-coded, with automatic fallbacks, streaming, and shared embedding generation across the platform.
- Built a three-stage **multi-agent GenAI pipeline** (Gatekeeper, Generator, Critic) on an async FastAPI microservice, automating candidate relevance scoring and personalised outreach in production.
- Shipped a recruiter copilot with **LLM tool-calling**: 6+ natural-language actions, token-by-token streaming, inline tool-result rendering, and an AI SDR voice-call step in the outreach workflow builder.
- Built the **role-based access control layer** for the product: per-role and per-project permissions, candidate-PII restrictions, audit trails, and permission changes propagating to live sessions over Socket.io.
- Owned the release delivery path: a Node.js document-export microservice, a branch-locked **zero-trust GitHub Actions pipeline** blocking unauthorized production deploys, health-checked Docker rollouts, and AWS infrastructure provisioned with Terraform and Ansible.
- Cut the dashboard's initial JavaScript payload **70% (10.7 MB → 3.3 MB)** via route-level code splitting, with remaining reductions sequenced in a documented performance plan.

HireQuotient | Frontend Engineer Intern

Bengaluru, India · Jul 2024 – Dec 2024

- Shipped customer-facing features across a large React dashboard (credit tracking, calendar sync with permission controls, email attachments, document export) and built early AI surfaces: chat-bot actions, AI-driven candidate filters, and AI summary templates.
- Migrated 200+ modules from Ant Design to MUI / Joy UI and custom components to keep design consistency.

PROJECTS

PolterTab | Browser Control for Any AI Assistant | github.com/abhinav162/poltertab | npmjs.com/package/poltertab

- Published an open-source **Model Context Protocol** server and Chrome extension that lets any MCP client drive an existing Chrome profile, keeping real logged-in sessions instead of launching a headless browser; **23 browser tools** over a stdio-to-WebSocket bridge, with **600+ npm downloads** in the first month.
- Handles the DOM cases that break scrapers: pierces open and closed **shadow roots**, searches child iframes in parallel, and falls back to **network interception** with scroll-triggered lazy loading for virtualized lists; record extraction reports fill rates that flag a wrong record boundary.

Synapse | Universal Memory Layer for AI Coding Tools | pypi.org/project/synapse-context-cli

- Built a 7-service platform that indexes codebases and serves context to Cursor, Claude, and Copilot over the **Model Context Protocol**, with a Next.js admin dashboard.
- Implemented vector search over **Qdrant** with a Redis caching layer and automated secret scanning that blocks credentials from being indexed; covered by 390+ pytest suites plus Vitest and Playwright specs.

EDUCATION

Lovely Professional University, B.Tech, Computer Science and Engineering, GPA 8.74 / 10.0

Completed Jul 2025